

D3: Query Answering Machine

Id	Category	Value	Description
1	Name	<i>Eric Gao</i>	I'm working alone
2	Andrew ID	<i>jingxia3</i>	My Andrew ID
3	Project Name	<i>Feeling Artsy</i>	Extracting human emotions with CV techniques

Below are some general things I've tried before you see what's in this document:

1. I initially used 4 classes. The accuracy are the same as random guessing. I merged image classes to increase the sample size for each class.
2. I initially had only 9 features. The training outcomes are not so ideal. I used polynomial features of them.
3. I initially had very small estimators for ensemble, resulting in poor testing accuracy for Random Forest. I made the estimator larger.
4. I initially had a big gap for CNN. I reduced batch size and increased epochs.

In this section, I chose the following models for my classification task:

- kNN
- Random Forest
- Convolutional Neural Network

D3.1: Process Iterations

I first changed the kernal to DataPreparation kernal to migrate the prepared dataset to here.

In [112...

df.info()

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 805 entries, 0 to 804
```

```
Data columns (total 71 columns):
```

#	Column	Non-Null Count	Dtype
0	ID	805 non-null	object
1	image	805 non-null	object
2	red	805 non-null	object
3	green	805 non-null	object
4	blue	805 non-null	object
5	rvar	805 non-null	float64
6	gvar	805 non-null	float64
7	bvar	805 non-null	float64
8	rskew	805 non-null	float64
9	gskew	805 non-null	float64
10	bskew	805 non-null	float64
11	rkurt	805 non-null	float64
12	gkurt	805 non-null	float64
13	bkurt	805 non-null	float64
14	l	805 non-null	float64
15	rvar	805 non-null	float64
16	gvar	805 non-null	float64
17	bvar	805 non-null	float64
18	rskew	805 non-null	float64
19	gskew	805 non-null	float64
20	bskew	805 non-null	float64
21	rkurt	805 non-null	float64
22	gkurt	805 non-null	float64
23	bkurt	805 non-null	float64
24	rvar^2	805 non-null	float64
25	rvar gvar	805 non-null	float64
26	rvar bvar	805 non-null	float64
27	rvar rskew	805 non-null	float64
28	rvar gskew	805 non-null	float64
29	rvar bskew	805 non-null	float64
30	rvar rkurt	805 non-null	float64
31	rvar gkurt	805 non-null	float64
32	rvar bkurt	805 non-null	float64
33	gvar^2	805 non-null	float64
34	gvar bvar	805 non-null	float64
35	gvar rskew	805 non-null	float64
36	gvar gskew	805 non-null	float64
37	gvar bskew	805 non-null	float64
38	gvar rkurt	805 non-null	float64
39	gvar gkurt	805 non-null	float64
40	gvar bkurt	805 non-null	float64
41	bvar^2	805 non-null	float64
42	bvar rskew	805 non-null	float64
43	bvar gskew	805 non-null	float64
44	bvar bskew	805 non-null	float64
45	bvar rkurt	805 non-null	float64
46	bvar gkurt	805 non-null	float64
47	bvar bkurt	805 non-null	float64
48	rskew^2	805 non-null	float64
49	rskew gskew	805 non-null	float64
50	rskew bskew	805 non-null	float64
51	rskew rkurt	805 non-null	float64
52	rskew gkurt	805 non-null	float64
53	rskew bkurt	805 non-null	float64
54	gskew^2	805 non-null	float64
55	gskew bskew	805 non-null	float64
56	gskew rkurt	805 non-null	float64
57	gskew gkurt	805 non-null	float64
58	gskew bkurt	805 non-null	float64

```

59 bskew^2      805 non-null    float64
60 bskew rkurt  805 non-null    float64
61 bskew gkurt  805 non-null    float64
62 bskew bkurt  805 non-null    float64
63 rkurt^2      805 non-null    float64
64 rkurt gkurt  805 non-null    float64
65 rkurt bkurt  805 non-null    float64
66 gkurt^2      805 non-null    float64
67 gkurt bkurt  805 non-null    float64
68 bkurt^2      805 non-null    float64
69 HOG          805 non-null    object
70 label        805 non-null    int64
dtypes: float64(64), int64(1), object(6)
memory usage: 446.6+ KB

```

```

In [113... hogdf = pd.DataFrame([np.hstack(i) for i in df['HOG']])
hogdf.shape, hogdf.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 805 entries, 0 to 804
Columns: 10000 entries, 0 to 9999
dtypes: float64(10000)
memory usage: 61.4 MB
((805, 10000), None)

```

Out[113]:

```

In [114... #import warnings filter
from warnings import simplefilter
# ignore all future warnings
simplefilter(action='ignore', category=FutureWarning)

```

Iteration 1: KNN

We humans often see emotions in images by relating them to similar images. For example, when looking at something with very bright colors, we might think it's happy because it relates to the fireworks we see.

As such, K Nearest Neighbors (KNN) model can be applied to perform this classification task. It is a non-parametric model that classifies data by considering its k closest neighbors. It falls under the supervised learning category, so our label can be useful.

Below I define some basic functions and import some libraries.

```

In [115... from sklearn.model_selection import cross_validate
from sklearn.neighbors import KNeighborsClassifier
from sklearn import metrics
def learn_kNN_classifier(features_data, targets, neighbors, voting):
    '''Set up a K-NN classifier and fits it to the given training data.
    Return: learned classifier.'''

    classifier = KNeighborsClassifier(n_neighbors=neighbors,
                                     weights=voting)

    classifier.fit(features_data, targets)

    return classifier

```

Below I define features and labels for training

```
In [116... features = pd.concat([df.iloc[:,5:69],hogdf],axis=1)
labels = df['label']
features.columns = features.columns.astype(str)
```

Below I use cross-validation to find the best k for kNN.

```
In [117... from sklearn import metrics

k_fold = 10

for k in [2,3,4,5,6,7,8,9]:

    classifier = learn_kNN_classifier(features,
                                      labels,
                                      k,
                                      'uniform')

    # if not specified either, K-Fold is the default CV method
    # and the cv argument specifies the number of foldings
    cv_results = cross_validate(classifier, features, labels,
                                cv=k_fold,
                                return_train_score=True)

    print('{}-NN] Mean test score: {:.3f} (std: {:.3f})'
          '\nMean train score: {:.3f} (std: {:.3f})\n'.format(k,
                                                              np.mean(cv_results['test_score']),
                                                              np.std(cv_results['test_score']),
                                                              np.mean(cv_results['train_score']),
                                                              np.std(cv_results['train_score'])))
```

```
[2-NN] Mean test score: 0.607 (std: 0.071)
Mean train score: 0.876 (std: 0.009)
```

```
[3-NN] Mean test score: 0.617 (std: 0.072)
Mean train score: 0.841 (std: 0.008)
```

```
[4-NN] Mean test score: 0.597 (std: 0.065)
Mean train score: 0.775 (std: 0.011)
```

```
[5-NN] Mean test score: 0.629 (std: 0.050)
Mean train score: 0.766 (std: 0.010)
```

```
[7-NN] Mean test score: 0.610 (std: 0.061)
Mean train score: 0.740 (std: 0.013)
```

```
[8-NN] Mean test score: 0.583 (std: 0.059)
Mean train score: 0.728 (std: 0.015)
```

```
[9-NN] Mean test score: 0.590 (std: 0.056)
Mean train score: 0.708 (std: 0.011)
```

The accuracies are a bit low. As k=5 has relatively the highest test and train scores, I would go with k=5 as the best one. Below is the confusion matrix.

```
In [118... classifier = learn_kNN_classifier(features,
                                    labels,
                                    5,
                                    'uniform')

Y_pred = classifier.predict(features)

from sklearn.metrics import confusion_matrix
```

```

import seaborn as sns

def show_confusion_matrix(true_labels, learned_labels, class_names):

    cmat = confusion_matrix(true_labels, learned_labels)

    plt.figure(figsize=(14, 5))

    plt.tick_params(labelsize=8)

    hm = sns.heatmap(cmat.T, square=True, annot=True, fmt='d', cbar=True,
                     xticklabels=class_names,
                     yticklabels=class_names,
                     cmap="seismic",
                     annot_kws={"size":12}, cbar_kws={'label': 'Counts'})

    # this is to set the last axis of the figure, the colorbar in this case
    hm.figure.axes[-1].yaxis.label.set_size(10) # fontsize of bar label
    hm.figure.axes[-1].tick_params(labelsize=8) # fontsize of ticks labels


    #plt.ylim(10, 0)

    plt.xlabel('True label', fontsize=9)
    plt.ylabel('Predicted label', fontsize=9)

    plt.show()

show_confusion_matrix(labels, Y_pred, [0, 1])

```



And because the low accuracy in both training and testing, this suggest that the model is too weak, and I'll considering using a stronger model.

Recall that there were around 350 images with class 0 and 450 images with class 1. It's expected that the model does better at recognizing class 1 images.

In sum, KNN seems to be doing a decent job with maximum accuracy around 0.67. Can we do better with another classifier?

Iteration 2: Random Forest

Let's see what Random Forest does for us.

As we see from KNN (which was a weak model for the task), it's better to use ensemble method to increase the model complexity in order to achieve better performance.

Random Forest is such an ensemble method that uses a bunch of decision tree models. For a decision tree, each internal node contains a criteria for subdividing samples from its parents. On its leaves, there would be labels that correspond to outcomes. In this case, we use a classification tree.

```
In [278... from sklearn.tree import DecisionTreeClassifier

k_fold = 10
max_samples = 0.3 # small sample size to reduce calculation
max_features = 0.7 # fewer features selected to increase diversity
n_estimators = 21 # that's how it's a forest

for depth in [5, 7, 9]:
    start = time.time()
    base = DecisionTreeClassifier(criterion='gini', max_depth=depth,
                                random_state=15288)
    classifier = BaggingClassifier(base, n_estimators=n_estimators,
                                max_samples=max_samples, max_features=max_features,
                                random_state=15288)

    cv_results = cross_validate(classifier, features, labels,
                                cv=k_fold,
                                return_train_score=True)

    print('[Random Forest of depth {}] Mean test score: {:.3f} (std: {:.3f})'
          '\nMean train score: {:.3f} (std: {:.3f})'.format(depth,
                                                            np.mean(cv_results['test_score']),
                                                            np.std(cv_results['test_score']),
                                                            np.mean(cv_results['train_score']),
                                                            np.std(cv_results['train_score'])))

    print('Time elapsed: {:.2f} seconds\n'.format(time.time()-start))
```

```
[Random Forest of depth 5] Mean test score: 0.614 (std: 0.047)
Mean train score: 0.880 (std: 0.014)
Time elapsed: 88.45 seconds
```


```
[Random Forest of depth 7] Mean test score: 0.647 (std: 0.036)
Mean train score: 0.894 (std: 0.009)
Time elapsed: 99.60 seconds
```

```
[Random Forest of depth 9] Mean test score: 0.624 (std: 0.042)
Mean train score: 0.899 (std: 0.006)
Time elapsed: 102.45 seconds
```

After some cross validation, max depth of 7 seems to be the most optimal. Let's see its confusion matrix.

```
In [279... base = DecisionTreeClassifier(criterion='gini', max_depth=7,
                                random_state=15288)
classifier = BaggingClassifier(base, n_estimators=n_estimators,
                              max_samples=max_samples, max_features=max_features,
                              random_state=15288)
classifier.fit(features, labels)
Y_pred = classifier.predict(features)

show_confusion_matrix(labels, Y_pred, [0, 1])
```



For Random Forest, the confusion matrix looks a lot better than the kNN. However, the generalization error is around the same as that of kNN. Considering Decision Tree to be a more complex model than kNN, there might be risks of overfitting.

Iteration 3: CNN

From the previous two attempts, we often find the model to be too weak to identify the subtle emotions in images.

Let's see if the all-powerful CNN create miracles. Convolutional Neural Network is commonly used for image related tasks, so I think it could be useful here.

```
In [128... import time
from time import process_time
import keras
from keras.models import Sequential
from keras.layers import Dense, Dropout, Activation, Flatten, Conv2D, MaxPooling2D
from keras.utils import np_utils
import tensorflow as tf
tf.random.set_seed(15288)
```

```
In [129... model = Sequential()
```

```
# input layer
model.add(Conv2D(128, kernel_size=(7, 7), input_shape=(100, 100, 3), activation='relu'))

# convolutional layer
model.add(Conv2D(64, kernel_size=(5, 5), activation='relu'))
model.add(Conv2D(64, kernel_size=(3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(3, 3)))
model.add(Dropout(0.2))

# flatten output of conv
model.add(Flatten())

# hidden layer
model.add(Dense(256, activation='relu'))
model.add(Dense(128, activation='relu'))
model.add(Dense(128, activation='relu'))
model.add(Dense(128, activation='relu'))
model.add(Dense(64, activation='relu'))
model.add(Dropout(0.2))

# output layer
model.add(Dense(2, activation='softmax'))
model.summary()
```

Model: "sequential_13"

Layer (type)	Output Shape	Param #
conv2d_51 (Conv2D)	(None, 94, 94, 128)	18944
conv2d_52 (Conv2D)	(None, 90, 90, 64)	204864
conv2d_53 (Conv2D)	(None, 88, 88, 64)	36928
max_pooling2d_17 (MaxPooling2D)	(None, 29, 29, 64)	0
dropout_31 (Dropout)	(None, 29, 29, 64)	0
flatten_13 (Flatten)	(None, 53824)	0
dense_73 (Dense)	(None, 256)	13779200
dense_74 (Dense)	(None, 128)	32896
dense_75 (Dense)	(None, 128)	16512
dense_76 (Dense)	(None, 128)	16512
dense_77 (Dense)	(None, 64)	8256
dropout_32 (Dropout)	(None, 64)	0
dense_78 (Dense)	(None, 2)	130
Total params: 14,114,242		
Trainable params: 14,114,242		
Non-trainable params: 0		

14 million parameters! Let's see how the model performs.

In [130...

```
model.compile(loss='categorical_crossentropy',
              optimizer='adam', metrics=['accuracy'])
```



```
In [131... from sklearn.model_selection import train_test_split
X_train, X_test, Y_train, Y_test = train_test_split(df['image'], df['label'],
                                                    test_size=0.2, random_state=1528)

img_shape = (100, 100, 3)

X_train = np.array([x.reshape(img_shape) for x in X_train])
X_test = np.array([x.reshape(img_shape) for x in X_test])
Y_train = np_utils.to_categorical(Y_train, 2)
Y_test = np_utils.to_categorical(Y_test, 2)
X_train.shape, X_test.shape, Y_train.shape, Y_test.shape
```

Out[131]: ((644, 100, 100, 3), (161, 100, 100, 3), (644, 2), (161, 2))

```
In [132... # Train
batch_size = 64
epochs = 20

start_time = time.process_time()

training = model.fit(X_train, Y_train,
                    batch_size=batch_size,
                    epochs=epochs,
                    validation_data=(X_test, Y_test),
                    verbose=2)

print('\n\nTraining time: {:.3f} sec\n'.format(time.process_time() - start_time))

score = model.evaluate(X_test, Y_test, verbose=0)

print('Test score:', score[0])
print('Test accuracy:', score[1])
```

Epoch 1/20
11/11 - 22s - loss: 20.7779 - accuracy: 0.4783 - val_loss: 0.8017 - val_accuracy: 0.4037 - 22s/epoch - 2s/step

Epoch 2/20
11/11 - 20s - loss: 0.7144 - accuracy: 0.5093 - val_loss: 0.6915 - val_accuracy: 0.5652 - 20s/epoch - 2s/step

Epoch 3/20
11/11 - 19s - loss: 0.6889 - accuracy: 0.5745 - val_loss: 0.6839 - val_accuracy: 0.5963 - 19s/epoch - 2s/step

Epoch 4/20
11/11 - 14s - loss: 0.6838 - accuracy: 0.5745 - val_loss: 0.6757 - val_accuracy: 0.5963 - 14s/epoch - 1s/step

Epoch 5/20
11/11 - 14s - loss: 0.7055 - accuracy: 0.5699 - val_loss: 0.6769 - val_accuracy: 0.5901 - 14s/epoch - 1s/step

Epoch 6/20
11/11 - 14s - loss: 0.6820 - accuracy: 0.5761 - val_loss: 0.6784 - val_accuracy: 0.5901 - 14s/epoch - 1s/step

Epoch 7/20
11/11 - 14s - loss: 0.6804 - accuracy: 0.5761 - val_loss: 0.6774 - val_accuracy: 0.5901 - 14s/epoch - 1s/step

Epoch 8/20
11/11 - 14s - loss: 0.6816 - accuracy: 0.5761 - val_loss: 0.6790 - val_accuracy: 0.5839 - 14s/epoch - 1s/step

Epoch 9/20
11/11 - 14s - loss: 0.6829 - accuracy: 0.5714 - val_loss: 0.6801 - val_accuracy: 0.5839 - 14s/epoch - 1s/step

Epoch 10/20
11/11 - 14s - loss: 0.6770 - accuracy: 0.5792 - val_loss: 0.6789 - val_accuracy: 0.5839 - 14s/epoch - 1s/step

Epoch 11/20
11/11 - 14s - loss: 0.6713 - accuracy: 0.5792 - val_loss: 0.6763 - val_accuracy: 0.5901 - 14s/epoch - 1s/step

Epoch 12/20
11/11 - 14s - loss: 0.6695 - accuracy: 0.5823 - val_loss: 0.6721 - val_accuracy: 0.5963 - 14s/epoch - 1s/step

Epoch 13/20
11/11 - 14s - loss: 0.6514 - accuracy: 0.5807 - val_loss: 0.7392 - val_accuracy: 0.5963 - 14s/epoch - 1s/step

Epoch 14/20
11/11 - 15s - loss: 0.6270 - accuracy: 0.5885 - val_loss: 0.6879 - val_accuracy: 0.5901 - 15s/epoch - 1s/step

Epoch 15/20
11/11 - 14s - loss: 0.5871 - accuracy: 0.6211 - val_loss: 0.7100 - val_accuracy: 0.5839 - 14s/epoch - 1s/step

Epoch 16/20
11/11 - 14s - loss: 0.5695 - accuracy: 0.6351 - val_loss: 0.8513 - val_accuracy: 0.5963 - 14s/epoch - 1s/step

Epoch 17/20
11/11 - 14s - loss: 0.5831 - accuracy: 0.6009 - val_loss: 0.7351 - val_accuracy: 0.5901 - 14s/epoch - 1s/step

Epoch 18/20
11/11 - 14s - loss: 0.5523 - accuracy: 0.6366 - val_loss: 0.8843 - val_accuracy: 0.6087 - 14s/epoch - 1s/step

Epoch 19/20
11/11 - 14s - loss: 0.5244 - accuracy: 0.6630 - val_loss: 2.1778 - val_accuracy: 0.6398 - 14s/epoch - 1s/step

Epoch 20/20
11/11 - 14s - loss: 0.5675 - accuracy: 0.6506 - val_loss: 0.7415 - val_accuracy: 0.6335 - 14s/epoch - 1s/step

Training time: 5477.969 sec

Test score: 0.7415297031402588

Test accuracy: 0.6335403919219971

In [133...

```
def show_learning_curves(learning_evolution, title=''):

    # list all data stored in the learning_evolution object
    print('Stored information:', [h for h in learning_evolution.history.keys()])

    # show evolution of accuracy
    plt.plot(learning_evolution.history['accuracy'])

    # show evolution of cross-validated accuracy
    if 'val_accuracy' in learning_evolution.history:
        plt.plot(learning_evolution.history['val_accuracy'])

    plt.title(title + ' Model accuracy')
    plt.ylabel('accuracy')
    plt.xlabel('epoch')

    plt.legend(['train', 'test'], loc='upper left')
    plt.show()


    # show evolution of loss
    plt.plot(learning_evolution.history['loss'])

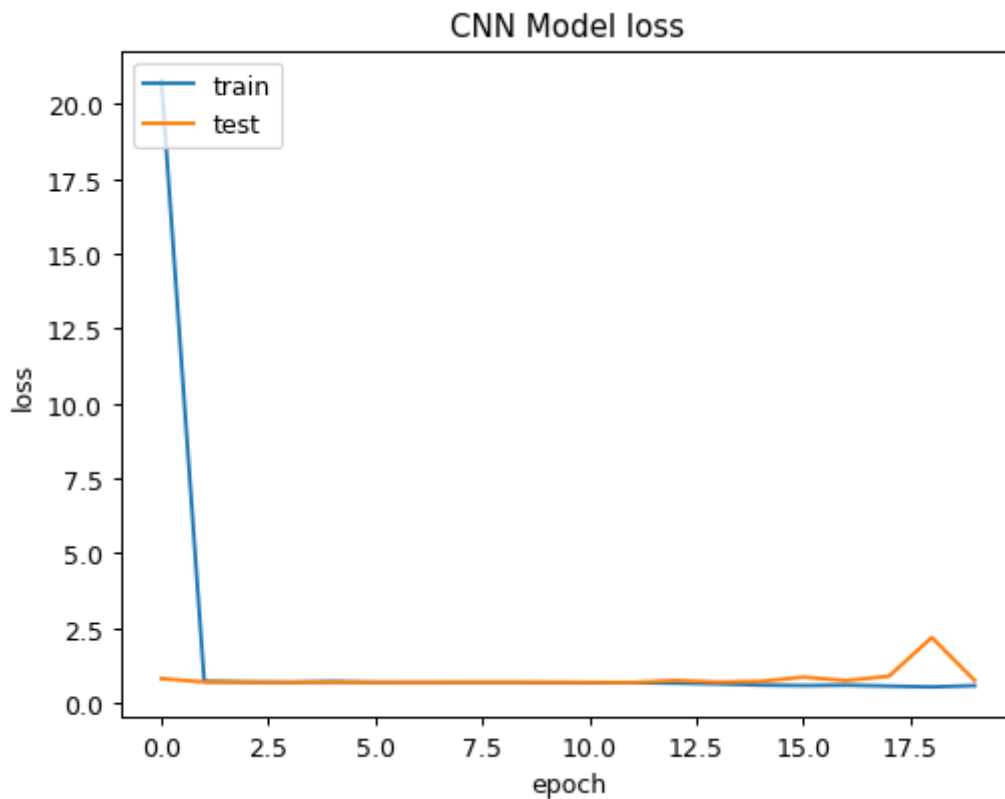
    # show evolution of cross-validated loss
    if 'val_loss' in learning_evolution.history:
        plt.plot(learning_evolution.history['val_loss'])

    plt.title(title + ' Model loss')
    plt.ylabel('loss')
    plt.xlabel('epoch')
    plt.legend(['train', 'test'], loc='upper left')
    plt.show()

show_learning_curves(training, 'CNN')
```

Stored information: ['loss', 'accuracy', 'val_loss', 'val_accuracy']





The model seems to have needed more training time. It might be a good idea to increase the number of epochs.

```
In [134... from keras.models import Model

def display_activation(activations, activation_index,
                      ncols = 5,
                      size=(12, 12), reshape=None):

    activation = activations[activation_index]

    print('Shape of layer: ', activation.shape)

    if len(activation.shape) > 2:
        total_fmaps = activation.shape[-1] # the last axis of the array
        nrows = int(np.ceil(total_fmaps / ncols))
    else:
        total_fmaps = 1
        nrows,ncols = (1,1)

    feature_map=0

    if (ncols, nrows) == (1,1):
        plt.yticks([], [])
        plt.xticks(fontsize=6)
        if reshape == None:
            plt.imshow(activation[0].reshape(-1,1).T) #cmap='gray')
        else:
            plt.imshow(activation[0].reshape(reshape).T)
    return

#fig, subplot = plt.subplots(nrows, ncols, figsize=size)
fig = plt.figure(figsize=size)
fig.subplots_adjust(hspace=0.23, wspace=0.22)

for row in range(nrows):
    for col in range(ncols):
```

```

subplot = fig.add_subplot(nrows, ncols, feature_map+1)
#subplot[row][col].imshow(activation[0, :, :, feature_map],
#                           cmap='gray')
#subplot[row][col].tick_params(labelsize = 3)
subplot.imshow(activation[0, :, :, feature_map])# cmap='gray')
subplot.tick_params(labelsize = 3)

feature_map += 1
if feature_map == total_fmmaps:
    return

```

In [135...

```

layer_outputs = [layer.output for layer in model.layers]

layer_names = [layer.name.split('_')[0] for layer in model.layers]

#print(layer_outputs)

activation_model = Model(inputs=model.input, outputs=layer_outputs)

img_index = 9

img_sample = X_train[img_index].reshape(1, *img_shape)

activations = activation_model.predict(img_sample)

for a, _ in enumerate(activations):
    print('Activation Layer [{}]: {} {}'.format(a, layer_names[a],
                                                  activations[a].shape))

```

```

1/1 [=====] - 0s 78ms/step
Activation Layer [0]: conv2d (1, 94, 94, 128)
Activation Layer [1]: conv2d (1, 90, 90, 64)
Activation Layer [2]: conv2d (1, 88, 88, 64)
Activation Layer [3]: max (1, 29, 29, 64)
Activation Layer [4]: dropout (1, 29, 29, 64)
Activation Layer [5]: flatten (1, 53824)
Activation Layer [6]: dense (1, 256)
Activation Layer [7]: dense (1, 128)
Activation Layer [8]: dense (1, 128)
Activation Layer [9]: dense (1, 128)
Activation Layer [10]: dense (1, 64)
Activation Layer [11]: dropout (1, 64)
Activation Layer [12]: dense (1, 2)

```

In [136...

```


print('Type of layer: ', layer_names[1])
display_activation(activations, 1, 8)

```

```

Type of layer: conv2d
Shape of layer: (1, 90, 90, 64)

```



Many of these feature maps don't have anything... Maybe we can reduce the model?

```
In [357... model = Sequential()

# input + convolutional layer
model.add(Conv2D(128, kernel_size=(7, 7), input_shape=(100, 100, 3), activation='relu'))
model.add(Dropout(0.2))
model.add(Conv2D(64, kernel_size=(5, 5), activation='relu'))
model.add(Dropout(0.2))
model.add(Conv2D(32, kernel_size=(3, 3), activation='relu'))
model.add(Dropout(0.2))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Conv2D(32, kernel_size=(7, 7), activation='relu'))
model.add(Dropout(0.2))
model.add(Conv2D(16, kernel_size=(5, 5), activation='relu'))
model.add(Dropout(0.2))
model.add(Conv2D(16, kernel_size=(3, 3), activation='relu'))
model.add(Dropout(0.2))
model.add(MaxPooling2D(pool_size=(2, 2)))

# flatten + dense + output layers
model.add(Flatten())
model.add(Dense(64, activation='relu'))
model.add(Dense(2, activation='softmax'))
model.summary()
```

Model: "sequential_85"

Layer (type)	Output Shape	Param #
conv2d_347 (Conv2D)	(None, 94, 94, 128)	18944
dropout_295 (Dropout)	(None, 94, 94, 128)	0
conv2d_348 (Conv2D)	(None, 90, 90, 64)	204864
dropout_296 (Dropout)	(None, 90, 90, 64)	0
conv2d_349 (Conv2D)	(None, 88, 88, 32)	18464
dropout_297 (Dropout)	(None, 88, 88, 32)	0
max_pooling2d_123 (MaxPooling2D)	(None, 44, 44, 32)	0
conv2d_350 (Conv2D)	(None, 38, 38, 32)	50208
dropout_298 (Dropout)	(None, 38, 38, 32)	0
conv2d_351 (Conv2D)	(None, 34, 34, 16)	12816
dropout_299 (Dropout)	(None, 34, 34, 16)	0
conv2d_352 (Conv2D)	(None, 32, 32, 16)	2320
dropout_300 (Dropout)	(None, 32, 32, 16)	0
max_pooling2d_124 (MaxPooling2D)	(None, 16, 16, 16)	0
flatten_81 (Flatten)	(None, 4096)	0
dense_330 (Dense)	(None, 64)	262208
dense_331 (Dense)	(None, 2)	130

=====

Total params: 569,954
Trainable params: 569,954
Non-trainable params: 0

=====

Much fewer parameters! Let's see how it performs.

In [358...

```
adam = tf.keras.optimizers.Adam(learning_rate=0.0005)
model.compile(loss='categorical_crossentropy',
               optimizer=adam, metrics=['accuracy'])
```

In [359...

```
from numpy.random import seed
seed(15288)
# Train
batch_size = 32
epochs = 50

start_time = time.process_time()

training = model.fit(X_train, Y_train,
                    batch_size=batch_size,
```

```
        epochs=epochs,  
        validation_data=(X_test,Y_test),  
        verbose=2)  
  
print('\n\nTraining time: {:.3f} sec\n'.format(time.process_time() - start_time))  
  
score = model.evaluate(X_test, Y_test, verbose=0)  
  
print('Test score:', score[0])  
print('Test accuracy:', score[1])
```


Epoch 1/50
21/21 - 16s - loss: 11.5805 - accuracy: 0.4907 - val_loss: 0.6843 - val_accuracy: 0.6273 - 16s/epoch - 752ms/step

Epoch 2/50
21/21 - 16s - loss: 0.7973 - accuracy: 0.4922 - val_loss: 0.6818 - val_accuracy: 0.6087 - 16s/epoch - 739ms/step

Epoch 3/50
21/21 - 16s - loss: 0.7170 - accuracy: 0.5683 - val_loss: 0.6696 - val_accuracy: 0.6398 - 16s/epoch - 754ms/step

Epoch 4/50
21/21 - 16s - loss: 0.7072 - accuracy: 0.5839 - val_loss: 0.6679 - val_accuracy: 0.6584 - 16s/epoch - 747ms/step

Epoch 5/50
21/21 - 16s - loss: 0.6754 - accuracy: 0.5963 - val_loss: 0.6755 - val_accuracy: 0.6025 - 16s/epoch - 747ms/step

Epoch 6/50
21/21 - 16s - loss: 0.6770 - accuracy: 0.5978 - val_loss: 0.6712 - val_accuracy: 0.6211 - 16s/epoch - 741ms/step

Epoch 7/50
21/21 - 16s - loss: 0.6715 - accuracy: 0.6320 - val_loss: 0.6660 - val_accuracy: 0.6211 - 16s/epoch - 755ms/step

Epoch 8/50
21/21 - 16s - loss: 0.6490 - accuracy: 0.6227 - val_loss: 0.6647 - val_accuracy: 0.6398 - 16s/epoch - 747ms/step

Epoch 9/50
21/21 - 16s - loss: 0.6416 - accuracy: 0.6444 - val_loss: 0.6572 - val_accuracy: 0.5963 - 16s/epoch - 748ms/step

Epoch 10/50
21/21 - 16s - loss: 0.6264 - accuracy: 0.6537 - val_loss: 0.6509 - val_accuracy: 0.6149 - 16s/epoch - 751ms/step

Epoch 11/50
21/21 - 16s - loss: 0.6373 - accuracy: 0.6413 - val_loss: 0.6615 - val_accuracy: 0.5963 - 16s/epoch - 751ms/step

Epoch 12/50
21/21 - 16s - loss: 0.6030 - accuracy: 0.6786 - val_loss: 0.6277 - val_accuracy: 0.6522 - 16s/epoch - 745ms/step

Epoch 13/50
21/21 - 16s - loss: 0.6453 - accuracy: 0.6429 - val_loss: 0.6559 - val_accuracy: 0.6087 - 16s/epoch - 750ms/step

Epoch 14/50
21/21 - 16s - loss: 0.6135 - accuracy: 0.6553 - val_loss: 0.6435 - val_accuracy: 0.6398 - 16s/epoch - 742ms/step

Epoch 15/50
21/21 - 16s - loss: 0.6507 - accuracy: 0.6087 - val_loss: 0.6492 - val_accuracy: 0.6273 - 16s/epoch - 753ms/step

Epoch 16/50
21/21 - 16s - loss: 0.6452 - accuracy: 0.6335 - val_loss: 0.6545 - val_accuracy: 0.6211 - 16s/epoch - 750ms/step

Epoch 17/50
21/21 - 16s - loss: 0.5832 - accuracy: 0.6879 - val_loss: 0.6917 - val_accuracy: 0.6025 - 16s/epoch - 751ms/step

Epoch 18/50
21/21 - 16s - loss: 0.6892 - accuracy: 0.6413 - val_loss: 0.6538 - val_accuracy: 0.6025 - 16s/epoch - 747ms/step

Epoch 19/50
21/21 - 16s - loss: 0.6391 - accuracy: 0.6568 - val_loss: 0.6571 - val_accuracy: 0.6087 - 16s/epoch - 758ms/step

Epoch 20/50
21/21 - 16s - loss: 0.6041 - accuracy: 0.6615 - val_loss: 0.6341 - val_accuracy: 0.6025 - 16s/epoch - 749ms/step

Epoch 21/50
21/21 - 16s - loss: 0.5691 - accuracy: 0.6755 - val_loss: 0.6277 - val_accuracy: 0.6025 - 16s/epoch - 777ms/step

Epoch 22/50

21/21 - 16s - loss: 0.5777 - accuracy: 0.6863 - val_loss: 0.6263 - val_accuracy: 0.6
460 - 16s/epoch - 764ms/step
Epoch 23/50
21/21 - 16s - loss: 0.5616 - accuracy: 0.7081 - val_loss: 0.5826 - val_accuracy: 0.6
335 - 16s/epoch - 747ms/step
Epoch 24/50
21/21 - 16s - loss: 0.5539 - accuracy: 0.7050 - val_loss: 0.5941 - val_accuracy: 0.6
335 - 16s/epoch - 751ms/step
Epoch 25/50
21/21 - 16s - loss: 0.5585 - accuracy: 0.7500 - val_loss: 0.6371 - val_accuracy: 0.5
963 - 16s/epoch - 748ms/step
Epoch 26/50
21/21 - 16s - loss: 0.5064 - accuracy: 0.7453 - val_loss: 0.5900 - val_accuracy: 0.6
770 - 16s/epoch - 758ms/step
Epoch 27/50
21/21 - 16s - loss: 0.5116 - accuracy: 0.7391 - val_loss: 0.5814 - val_accuracy: 0.6
832 - 16s/epoch - 745ms/step
Epoch 28/50
21/21 - 16s - loss: 0.4659 - accuracy: 0.7609 - val_loss: 0.6272 - val_accuracy: 0.6
584 - 16s/epoch - 742ms/step
Epoch 29/50
21/21 - 16s - loss: 0.5259 - accuracy: 0.7360 - val_loss: 0.5490 - val_accuracy: 0.6
832 - 16s/epoch - 746ms/step
Epoch 30/50
21/21 - 16s - loss: 0.4758 - accuracy: 0.7702 - val_loss: 0.5421 - val_accuracy: 0.7
019 - 16s/epoch - 749ms/step
Epoch 31/50
21/21 - 16s - loss: 0.4239 - accuracy: 0.7873 - val_loss: 0.5602 - val_accuracy: 0.6
957 - 16s/epoch - 742ms/step
Epoch 32/50
21/21 - 16s - loss: 0.3873 - accuracy: 0.7857 - val_loss: 0.6235 - val_accuracy: 0.6
708 - 16s/epoch - 749ms/step
Epoch 33/50
21/21 - 16s - loss: 0.4003 - accuracy: 0.7780 - val_loss: 0.5872 - val_accuracy: 0.7
267 - 16s/epoch - 742ms/step
Epoch 34/50
21/21 - 16s - loss: 0.3681 - accuracy: 0.8043 - val_loss: 0.6274 - val_accuracy: 0.6
957 - 16s/epoch - 747ms/step
Epoch 35/50
21/21 - 16s - loss: 0.3512 - accuracy: 0.8292 - val_loss: 0.6185 - val_accuracy: 0.6
522 - 16s/epoch - 743ms/step
Epoch 36/50
21/21 - 16s - loss: 0.3842 - accuracy: 0.8059 - val_loss: 0.5662 - val_accuracy: 0.7
329 - 16s/epoch - 745ms/step
Epoch 37/50
21/21 - 16s - loss: 0.3715 - accuracy: 0.8106 - val_loss: 0.6782 - val_accuracy: 0.6
832 - 16s/epoch - 746ms/step
Epoch 38/50
21/21 - 16s - loss: 0.3418 - accuracy: 0.8261 - val_loss: 0.6930 - val_accuracy: 0.6
584 - 16s/epoch - 744ms/step
Epoch 39/50
21/21 - 16s - loss: 0.3546 - accuracy: 0.8307 - val_loss: 0.7460 - val_accuracy: 0.6
584 - 16s/epoch - 743ms/step
Epoch 40/50
21/21 - 16s - loss: 0.3193 - accuracy: 0.8339 - val_loss: 0.7596 - val_accuracy: 0.6
957 - 16s/epoch - 744ms/step
Epoch 41/50
21/21 - 16s - loss: 0.2700 - accuracy: 0.8680 - val_loss: 0.9553 - val_accuracy: 0.6
708 - 16s/epoch - 747ms/step
Epoch 42/50
21/21 - 16s - loss: 0.4999 - accuracy: 0.7655 - val_loss: 0.6841 - val_accuracy: 0.6
460 - 16s/epoch - 742ms/step
Epoch 43/50
21/21 - 16s - loss: 0.3984 - accuracy: 0.7997 - val_loss: 0.5775 - val_accuracy: 0.7

```

081 - 16s/epoch - 746ms/step
Epoch 44/50
21/21 - 16s - loss: 0.3138 - accuracy: 0.8385 - val_loss: 0.5805 - val_accuracy: 0.7
205 - 16s/epoch - 753ms/step
Epoch 45/50
21/21 - 16s - loss: 0.2634 - accuracy: 0.8665 - val_loss: 0.6384 - val_accuracy: 0.7
143 - 16s/epoch - 753ms/step
Epoch 46/50
21/21 - 16s - loss: 0.2396 - accuracy: 0.8758 - val_loss: 0.7045 - val_accuracy: 0.6
894 - 16s/epoch - 747ms/step
Epoch 47/50
21/21 - 16s - loss: 0.2257 - accuracy: 0.8944 - val_loss: 0.8371 - val_accuracy: 0.7
019 - 16s/epoch - 767ms/step
Epoch 48/50
21/21 - 16s - loss: 0.2460 - accuracy: 0.8913 - val_loss: 0.6724 - val_accuracy: 0.7
019 - 16s/epoch - 756ms/step
Epoch 49/50
21/21 - 16s - loss: 0.2096 - accuracy: 0.8882 - val_loss: 1.2223 - val_accuracy: 0.7
081 - 16s/epoch - 751ms/step
Epoch 50/50
21/21 - 16s - loss: 0.2503 - accuracy: 0.8991 - val_loss: 0.5224 - val_accuracy: 0.7
578 - 16s/epoch - 743ms/step

```

Training time: 14056.797 sec


Test score: 0.5224094986915588


Test accuracy: 0.7577639818191528

Wow! Testing accuracy is over 75%! CNN is magic indeed!

In [360... `show_learning_curves(training, 'CNN')`

Stored information: ['loss', 'accuracy', 'val_loss', 'val_accuracy']





The training curves look a bit stepwise? After 10 epochs, both training and testing accuracy begins to stabilize, but the testing error starts to increase. There might be an overfitting happening. However, it is reasonable to say that CNN performed around 10% better than random forest.

Let's see its performance on training data!

```
In [361... from sklearn.metrics import confusion_matrix
import seaborn as sns

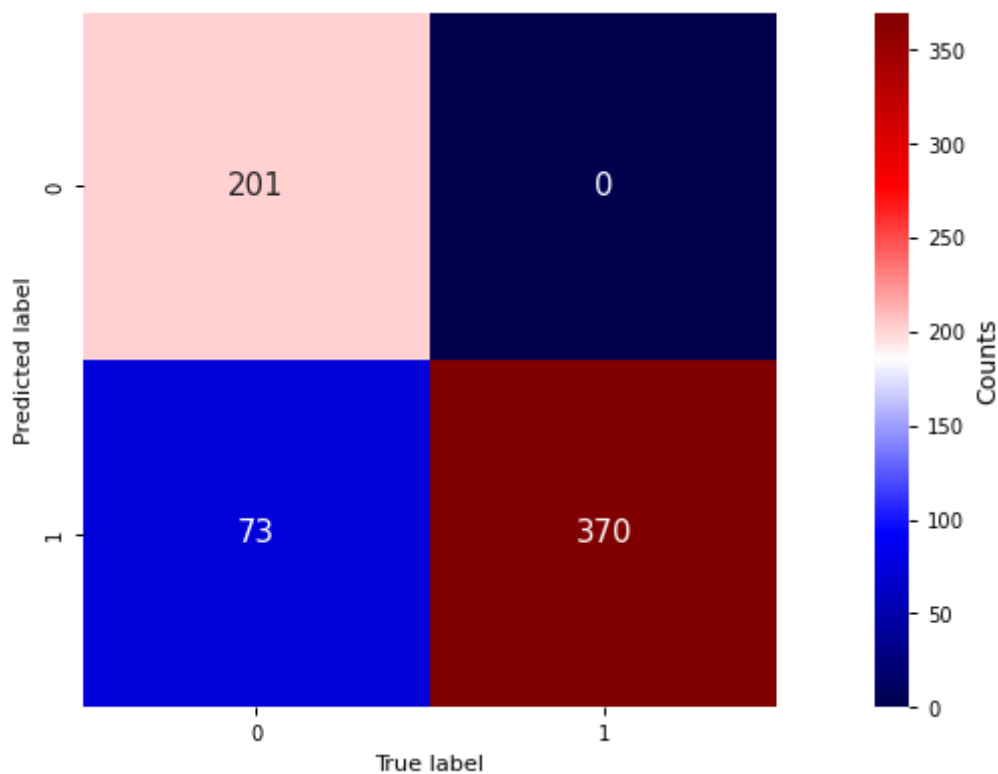
Y_prediction = model.predict(X_train)

# Convert predictions classes to one hot vectors
Y_pred_classes = np.argmax(Y_prediction, axis = 1)

# Convert test data to one hot vectors
Y_true = np.argmax(Y_train, axis = 1)

show_confusion_matrix(Y_true, Y_pred_classes, [0, 1])

21/21 [=====] - 3s 134ms/step
```



How about for testing data?

In [362...

```
from sklearn.metrics import confusion_matrix
import seaborn as sns


Y_prediction = model.predict(X_test)

# Convert predictions classes to one hot vectors
Y_pred_classes = np.argmax(Y_prediction, axis = 1)

# Convert test data to one hot vectors
Y_true = np.argmax(Y_test, axis = 1)

show_confusion_matrix(Y_true, Y_pred_classes, [0, 1])
```

6/6 [=====] - 1s 116ms/step



Not bad! We see that it rarely misclassify negative images as positive images, but sometimes misclassify positive images as negative ones.

Final words

Id	Model	Parameters	Training accuracy	Testing accuracy
1	kNN	neighbors=5, weights='distance'	0.740	0.629
2	Random Forest	max_depth=7, criterion='gini', n_estimators = 21	0.894	0.647
3	CNN	(see model summary above, epoch 49)	0.899	0.758

Clearly, CNN is the most optimal model for our task. Let's use it in the app!

In []: